

Project 2 Preliminary Design

General Plan of Attack

For this project, I will follow a development style known as test-driven development. As the name suggests, I will initially write unit tests for functionality I intend to have or add into the kernel. Unit tests will correspond with the required system calls.

Locking

I plan to use reader-writer locks for all system calls, where reset, block, and unblock system calls will be writers, and query system calls will be readers.

Storing Data

One of the important characteristics of these systems is that their data will be queried every time a socket or file is created or accessed.

Firewall

For managing and storing firewall entries, I plan to keep a red-black tree of custom structs where the key of each node is the port number. I will use the Linux provided implementation in **include/linux/rbtree.h**. Red-black trees seem to make the most sense here because port numbers can easily be used here as node keys, so implementing various functions for the tree will not be overly complex. Although a hash table could improve upon the $O(\lg n)$ runtime of various operations, depending on the size of the table, the amount of kernel memory in use at any given moment would be fairly large considering the 2^{16} ports available to block.

File Access Control

For managing and storing blocked files, I plan to keep a (chaining) hash table of custom structs where the file path (which starts from /) will be used as input to the hash function. I will use the Linux provided implementation in **include/linux/hashtable.h**. I decided on a hash table, after comparing them with B+ trees and radix trees (Patricia tries), because of their inherent ability to categorize various types of data (i.e. string representations of file paths) and near-constant operation times.

Kernel Modifications

Below, I have traced out various system calls that I believe to be relevant to each of the project's requirements. The legend is as follows:

Legend

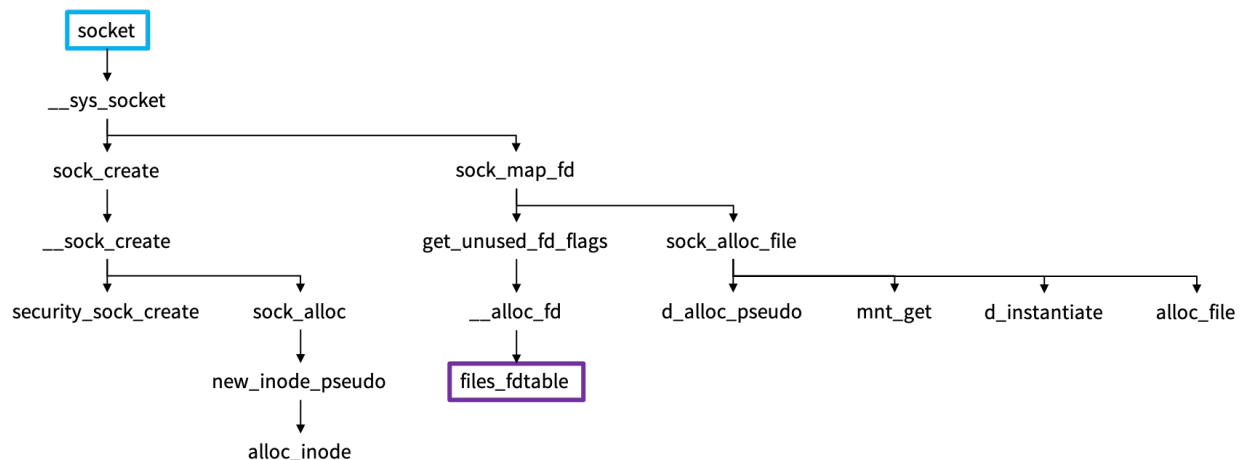
System call

Macro

Function of interest

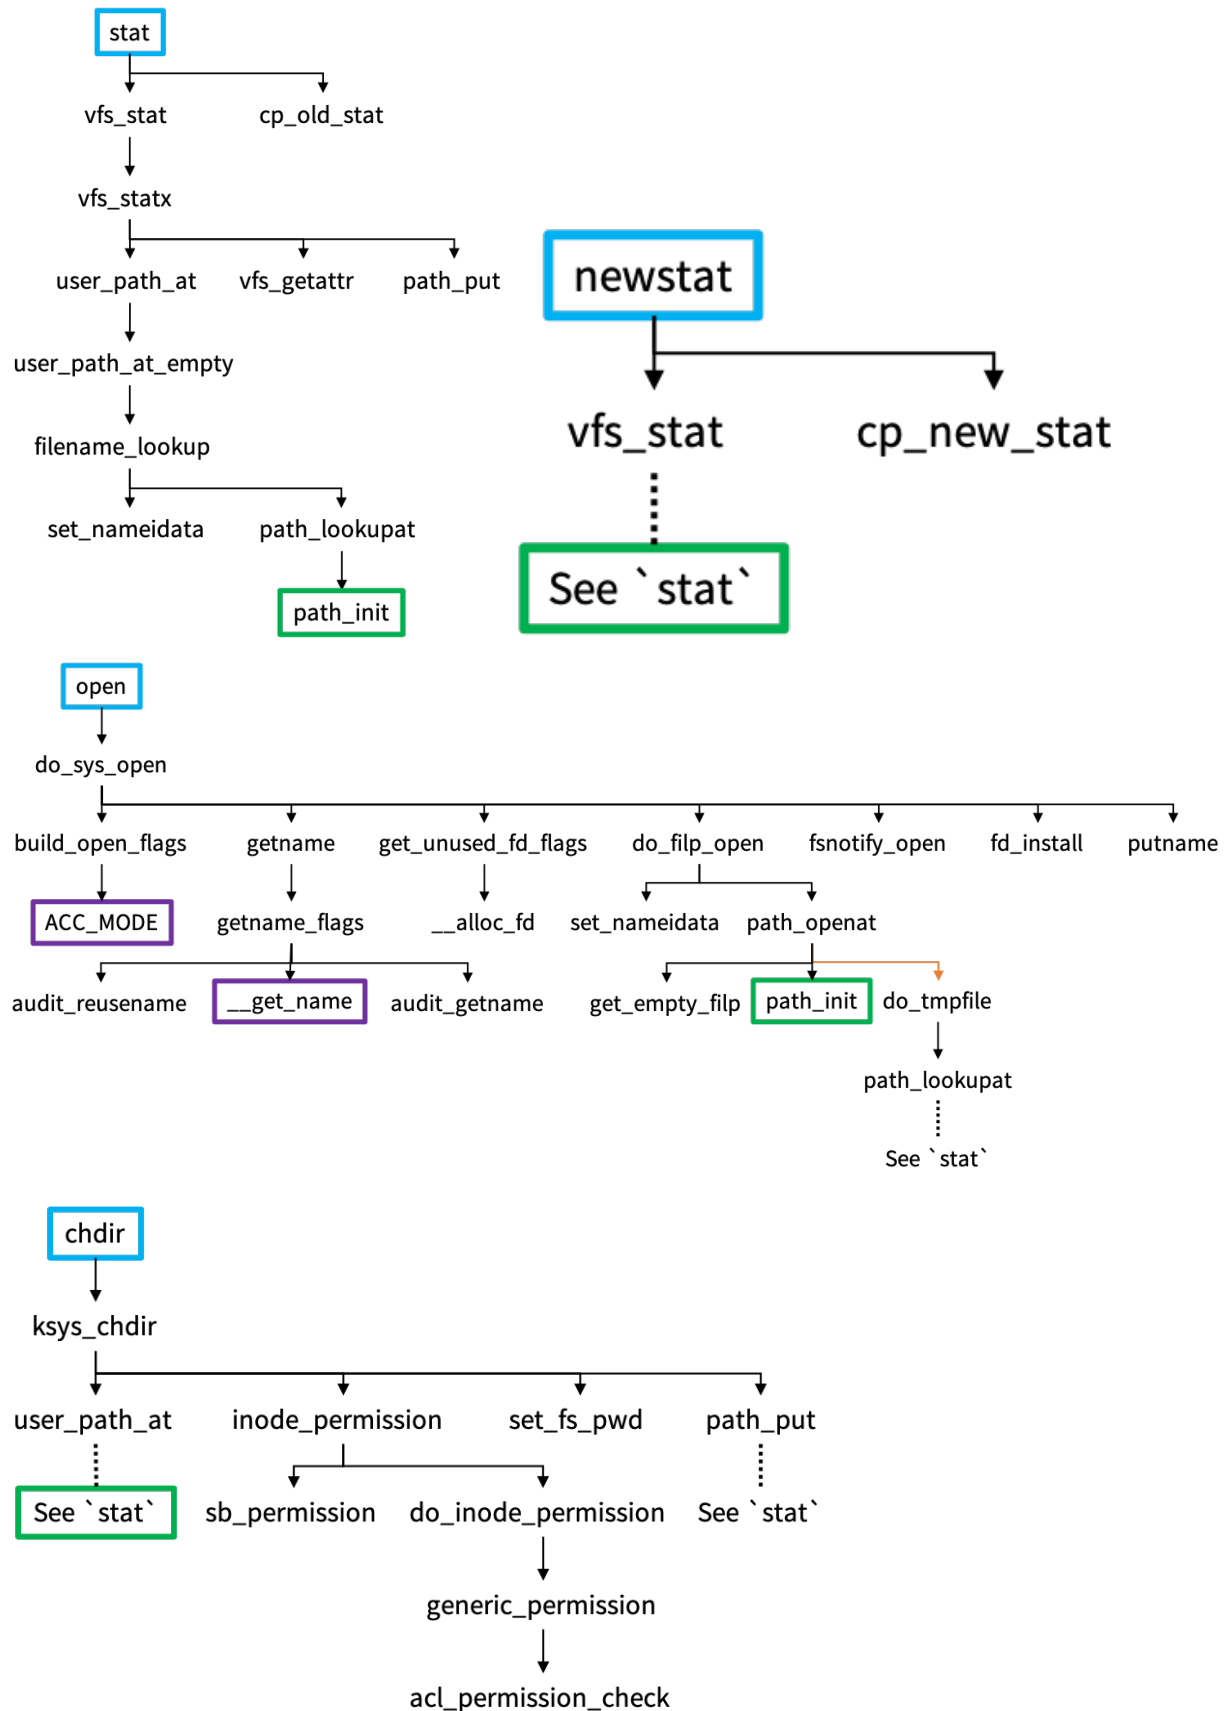
Firewall

The socket system call (in **net/socket.c**) is of most interest for the firewall, since it handles the creation of a socket, the very action that would need to be blocked.



File Access Control

The function of interest for file access control is `path_init` in **fs/namei.c**. This function appears as a place of convergence for many system calls that deal with the filesystem.



Conclusion

Please, any advice you are at liberty to share with me will be helpful. It took me multiple hours to simply realize I was looking at the wrong version of kernel source code in the Linux Cross Reference, let alone understanding how they work and spending time researching data structures in the kernel and how to use them. Thank you.

Sources

The following sources were useful in the design process and for understanding various kernel functionality:

- Sockets and Networking
 - https://www.tutorialspoint.com/unix_sockets/index.htm
 - <http://man7.org/linux/man-pages/man2/bind.2.html>
- Filesystem
 - <https://www.howtoforge.com/linux-stat-command/>
 - http://www.tutorialspoint.com/unix_system_calls/chdir.htm
 - <https://eklitzke.org/path-max-is-tricky>
- Kernel Data Structures
 - <http://luisbg.blogalia.com/historias/74062>
 - https://en.wikipedia.org/wiki/Radix_tree
 - <http://www.voidcn.com/article/p-fjdzcnrj-boo.html>
 - <https://lwn.net/Articles/175432/>
 - <http://events17.linuxfoundation.org/sites/events/files/slides/LinuxConNA2016%20-%20Radix%20Tree.pdf>
 - <https://kernelnewbies.org/FAQ/Hashtables>
- Miscellaneous
 - <https://kernelnewbies.org/FAQ/LikelyUnlikely>